

# Part3\_Model

April 6, 2021

## 1 File Information

Name: Amie Davis

Course: DSC680 - Applied Data Science

Project: College Recommendation Engine

Purpose: Build model(s)

Usage: Python 3.7.6

Developed using Jupyter Notebook 6.0.3

## 2 Data Source

College Scorecard is managed by the US Department of Education and provides results such as costs and graduation rates.

College Scorecard. (n.d.). Retrieved March 15, 2021, from <https://collegescorecard.ed.gov/>

## 3 References

Albon, C. (2018). Machine learning with Python cookbook: practical solutions from preprocessing to deep learning. O'Reilly.

<https://towardsdatascience.com/build-your-own-clustering-based-recommendation-engine-in-15-minutes-bddddd591d394>

## 4 Part 3

In Part 3, I will build a k-Means clustering model to perform unsupervised learning to group similar colleges.

### 4.1 Import required packages

```
[1]: # Suppress Warnings
      #import warnings
      #warnings.filterwarnings('ignore')

      import pandas as pd
```

```
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import pickle
```

## 5 Prepare Data

```
[2]: # Load data into dataframe
data_file = "Data\Scorecard\Cleaned_Scorecard.csv"
df = pd.read_csv(data_file)
```

### 5.1 Eliminate non-numeric features

```
[3]: print(df.dtypes)
```

```
Unnamed: 0          int64
UNITID             int64
SAT_AVG            int64
ACT_MEDIAN         int64
TUITION_OUT_ST     int64
LOCALE             int64
LATITUDE           float64
LONGITUDE          float64
ADM_RATE_ALL       float64
SIZE               int64
ONLINE_ONLY        int64
Y4_COMPLETION_RT   float64
PT_RETENTION_RT    float64
UG_INCOMP_1Y_REPAY_RT float64
UG_GRAD_4Y_REPAY_RT float64
UG_PLUS_1Y_REPAY_RT float64
Y4_COMPLETION_RT_POOLED float64
FED_LN_AWD_RT      float64
Y5_UG_REPAY_RT     float64
Y5_PLUS_LN_REPAY_RT float64
DBRR10_FED_UG_RT   float64
DBRR10_PP_UG_RT    float64
OPEID              int64
OPEID6             int64
SCHOOL_NAME        object
CITY                object
STABBR             object
ZIP                object
MAIN_CAMPUS        int64
NUM_BRANCH          int64
OWNERSHIP           int64
```

```

ST_FIPS                int64
REGION                 int64
NPT41_PRIV             float64
PVT_INCOME_30_40K      float64
TUITION_IN_ST          int64
INSTR_EXP_PER_FTE       int64
Y3_LN_DEFAULT_RT       float64
MEDIAN_DEBT            float64
MEDIAN_PLUS_LN_DEBT     int64
MEDIAN_PLUS_LN_DEBT_GRADS float64
Y2_LN_DEFAULT_RT       float64
Y2_LN_DELINQ_RT        float64
FED_UG_PAIDINFULL      float64
UG_GRAD_2Y_FORBEAR_RT  float64
LPSTAFFORD_CNT         float64
NUM_STU_PLUS_LN_BAL     int64
LPPPLUS_AMT            float64
FED_SCHOOL_CD          object
dtype: object

```

```

[4]: # Keep only numeric fields
train_df = df.select_dtypes(include=[np.number])

train_df.head()

```

```

[4]:   Unnamed: 0  UNITID  SAT_AVG  ACT_MEDIAN  TUITION_OUT_ST  LOCALE  LATITUDE  \
0           1  100654     957           18         18354      12  34.783368
1           2  100663    1220           25         19704      12  33.505697
2           3  100690      0           0          6900      12  32.362609
3           4  100706    1314           28         22362      12  34.724557
4           5  100724     972           18         19396      12  32.364317

```

```

      LONGITUDE  ADM_RATE_ALL  SIZE  ...  MEDIAN_DEBT  MEDIAN_PLUS_LN_DEBT  \
0 -86.568502      0.8986  4990  ...      3606.0         14838
1 -86.799345      0.9211  13186 ...      7504.0         16145
2 -86.174010      1.0000   351  ...       514.0           0
3 -86.640449      0.8087  7458  ...      3021.0        13524
4 -86.295677      0.9774  3903  ...      3609.0        15351

```

```

      MEDIAN_PLUS_LN_DEBT_GRADS  Y2_LN_DEFAULT_RT  Y2_LN_DELINQ_RT  \
0          16106.0          0.172640          0.074776
1          16954.0          0.060389          0.034483
2              NaN          0.157191          0.050167
3          16550.0          0.055035          0.031390
4          18952.0          0.173804          0.061713

```

```

      FED_UG_PAIDINFULL  UG_GRAD_2Y_FORBEAR_RT  LPSTAFFORD_CNT  \

```

0	0.007926	0.266409	31374.0
1	0.059675	0.156121	56997.0
2	NaN	0.000000	4463.0
3	0.092947	0.100349	19702.0
4	0.007872	0.286976	34246.0

	NUM_STU_PLUS_LN_BAL	LPPPLUS_AMT
0	5201	113949554.0
1	3727	87421879.0
2	16	309618.0
3	1397	29789762.0
4	4602	95359062.0

[5 rows x 44 columns]

```
[5]: # Find null records
count_nan_in_df = train_df.isnull().sum()
print (count_nan_in_df)
```

Unnamed: 0	0
UNITID	0
SAT_AVG	0
ACT_MEDIAN	0
TUITION_OUT_ST	0
LOCALE	0
LATITUDE	0
LONGITUDE	0
ADM_RATE_ALL	0
SIZE	0
ONLINE_ONLY	0
Y4_COMPLETION_RT	0
PT_RETENTION_RT	0
UG_INCOMP_1Y_REPAY_RT	0
UG_GRAD_4Y_REPAY_RT	0
UG_PLUS_1Y_REPAY_RT	0
Y4_COMPLETION_RT_POOLED	0
FED_LN_AWD_RT	0
Y5_UG_REPAY_RT	0
Y5_PLUS_LN_REPAY_RT	0
DBRR10_FED_UG_RT	207
DBRR10_PP_UG_RT	362
OPEID	0
OPEID6	0
MAIN_CAMPUS	0
NUM_BRANCH	0
OWNERSHIP	0
ST_FIPS	0
REGION	0

NPT41_PRIV	226
PVT_INCOME_30_40K	758
TUITION_IN_ST	0
INSTR_EXP_PER_FTE	0
Y3_LN_DEFAULT_RT	0
MEDIAN_DEBT	129
MEDIAN_PLUS_LN_DEBT	0
MEDIAN_PLUS_LN_DEBT_GRADS	407
Y2_LN_DEFAULT_RT	0
Y2_LN_DELIQ_RT	0
FED_UG_PAIDINFULL	427
UG_GRAD_2Y_FORBEAR_RT	0
LPSTAFFORD_CNT	89
NUM_STU_PLUS_LN_BAL	0
LPPPLUS_AMT	211

dtype: int64

[6]: *# Handle null Values*

```
# Set null rate data to 0
train_df["DBRR10_FED_UG_RT"].fillna(0, inplace = True)
train_df["DBRR10_PP_UG_RT"].fillna(0, inplace = True)

# Set null monetary data to 0
train_df["FED_UG_PAIDINFULL"].fillna(0, inplace = True)
train_df["LPPPLUS_AMT"].fillna(0, inplace = True)
train_df["PVT_INCOME_30_40K"].fillna(0, inplace = True)

# Impute null median data to median of values in other records
train_df['MEDIAN_DEBT'].fillna(train_df['MEDIAN_DEBT'].median(), inplace=True)
train_df['MEDIAN_PLUS_LN_DEBT_GRADS'].
    →fillna(train_df['MEDIAN_PLUS_LN_DEBT_GRADS'].median(), inplace=True)

# Some schools have no students with reported income < $30
# Set cost to 0 from NA
train_df["NPT41_PRIV"].fillna(0, inplace = True)
train_df["LPSTAFFORD_CNT"].fillna(0, inplace = True)
```

C:\Users\amomu\Anaconda3\lib\site-packages\pandas\core\series.py:4469:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
downcast=downcast,

[7]: *# Check null records*

```
count_nan_in_df = train_df.isnull().sum()
```

```
print (count_nan_in_df)
```

Unnamed: 0	0
UNITID	0
SAT_AVG	0
ACT_MEDIAN	0
TUITION_OUT_ST	0
LOCALE	0
LATITUDE	0
LONGITUDE	0
ADM_RATE_ALL	0
SIZE	0
ONLINE_ONLY	0
Y4_COMPLETION_RT	0
PT_RETENTION_RT	0
UG_INCOMP_1Y_REPAY_RT	0
UG_GRAD_4Y_REPAY_RT	0
UG_PLUS_1Y_REPAY_RT	0
Y4_COMPLETION_RT_POOLED	0
FED_LN_AWD_RT	0
Y5_UG_REPAY_RT	0
Y5_PLUS_LN_REPAY_RT	0
DBRR10_FED_UG_RT	0
DBRR10_PP_UG_RT	0
OPEID	0
OPEID6	0
MAIN_CAMPUS	0
NUM_BRANCH	0
OWNERSHIP	0
ST_FIPS	0
REGION	0
NPT41_PRIV	0
PVT_INCOME_30_40K	0
TUITION_IN_ST	0
INSTR_EXP_PER_FTE	0
Y3_LN_DEFAULT_RT	0
MEDIAN_DEBT	0
MEDIAN_PLUS_LN_DEBT	0
MEDIAN_PLUS_LN_DEBT_GRADS	0
Y2_LN_DEFAULT_RT	0
Y2_LN_DELINQ_RT	0
FED_UG_PAIDINFULL	0
UG_GRAD_2Y_FORBEAR_RT	0
LPSTAFFORD_CNT	0
NUM_STU_PLUS_LN_BAL	0
LPPPLUS_AMT	0

dtype: int64

## 5.2 Standardization

```
[8]: # Standardize features into array
scaler = StandardScaler()
features_std = scaler.fit_transform(train_df)

features_std

[8]: array([[ -1.73120939, -1.1602827 ,  0.45343115, ..., -0.08513354,
          0.70559181,  0.30577287],
          [-1.72952616, -1.16019784,  0.91757109, ...,  0.08830825,
          0.35705198,  0.13645898],
          [-1.72784292, -1.15994327, -1.23547352, ..., -0.26729379,
          -0.5204455 , -0.41953831],
          ...,
          [ 1.72784292,  2.53671497, -1.23547352, ...,  0.37345189,
          0.03097572, -0.06963586],
          [ 1.72952616,  2.53711097, -1.23547352, ..., -0.19134574,
          -0.35705267, -0.29983386],
          [ 1.73120939,  2.53862898, -1.23547352, ..., -0.11424696,
          -0.27121824, -0.27673791]])
```

## 6 Model Creation and Evaluation

### 6.1 Build Model

```
[9]: # Use Elbow Method to determine best value for k

# Calculates Within-Cluster-Sum of Squared Errors for k from 1 to kmax
def calculate_WSS(points, kmax):

    # Initiate list
    sse = []

    for k in range(1, kmax+1):
        kmeans = KMeans(n_clusters = k).fit(points)
        centroids = kmeans.cluster_centers_
        pred_clusters = kmeans.predict(points)
        curr_sse = 0

        # Calculate square of Euclidean distance of each point from its cluster_
        ↪center and add to current WSS
        for i in range(len(points)):
            curr_center = centroids[pred_clusters[i]]
            curr_sse += (points[i, 0] - curr_center[0]) ** 2 + (points[i, 1] -
            ↪curr_center[1]) ** 2
```

```

    # Return pairs of k and WWS so they can be plotted
    sse_tup = (k, curr_sse)
    sse.append(sse_tup)

    return sse

# Test output
calculate_WSS(features_std, 10)

```

```

[9]: [(1, 4115.9999999999996),
      (2, 3059.4212503151994),
      (3, 3644.599246722471),
      (4, 2845.5805426818197),
      (5, 2935.4186427435056),
      (6, 3043.9849555606074),
      (7, 2882.408960006184),
      (8, 2849.207259898587),
      (9, 2784.8486009196317),
      (10, 2766.7883661519377)]

```

```

[31]: # Plot Elbow Curve (WSS vs. k)
#fig, ax = plt.subplots() # Create a figure containing a single axes.
# Setup figure w/o axes lines
plt.figure()
plt.axes(frameon=False)

# Plot data
plt.plot(calculate_WSS(features_std, 100), color='c', linewidth=2.0)


# Set plot parameters
plt.title('Sum of Squared Errors', loc='left', fontsize=18)
plt.xlabel('k')
plt.ylabel('SSE')

# Save plot to file
plt.savefig('Elbow_Curve.png', bbox_inches='tight')

# Show Plot
plt.show()

```





```
[11]: # Create k-Means model
      # Selecting k=30 based on the
      cluster = KMeans(n_clusters=30, random_state=0, n_jobs=-1)
```

```
[12]: # Train model
      kmeans = cluster.fit(features_std)
```

C:\Users\amomu\Anaconda3\lib\site-packages\sklearn\cluster\\_kmeans.py:793:  
FutureWarning: 'n\_jobs' was deprecated in version 0.23 and will be removed in  
1.0 (renaming of 0.25).  
" removed in 1.0 (renaming of 0.25).", FutureWarning)

```
[13]: # Save machine learning model
      filename = 'kmeans_model.sav'
      pickle.dump(kmeans, open(filename, 'wb'))
```

### 6.1.1 Model Evaluation

```
[14]: # View centroid values
      print(kmeans.cluster_centers_)
```

```
[[ 1.51593222  2.2711597 -1.23547352 ...  9.37611862  1.51617153
    0.32900296]
 [-0.87711334 -0.64054999  0.51630695 ... -0.23810965 -0.34628419
```

```

-0.30935792]
[ 0.52452695 -0.02651169  0.52852398 ... -0.239954   -0.30940496
 -0.27394961]
...
[ 1.26425136  1.25232638 -1.23547352 ... -0.28551027 -0.52422883
 -0.42151446]
[-0.2491191  -0.08656262 -1.23547352 ... -0.26012566 -0.48678957
 -0.39182638]
[-0.45663128 -0.47645366 -1.17628204 ... -0.18575183 -0.39311939
 -0.3611353  ]

```

```

[15]: # View cluster label
kmeans.labels_

```

```

[15]: array([21,  5, 10, ..., 19, 14, 21])

```

## 7 Model Deployment

```

[16]: # Function gets the cluster of a school
def cluster_predict(unit_id):

    # Find record in clean numeric df
    record_df = train_df.loc[train_df['UNITID'] == unit_id]

    # Standardize into array
    std_array = scaler.fit_transform(record_df)

    # Assign cluster for specific school
    prediction = kmeans.predict(std_array)
    return prediction

# Test function
# print(cluster_predict(100654))

```

```

[17]: # Store cluster labels for each school into original df

# Create new column
df['CLUSTERLABEL'] = ""

# Assign cluster label for each school
df['CLUSTERLABEL']=df.apply(lambda x: cluster_predict(df['UNITID']), axis=0)

df.head()

```

```

[17]:   Unnamed: 0  UNITID  SAT_AVG  ACT_MEDIAN  TUITION_OUT_ST  LOCALE  LATITUDE  \
0           1  100654      957           18           18354      12  34.783368
1           2  100663     1220           25           19704      12  33.505697

```

2	3	100690	0	0	6900	12	32.362609
3	4	100706	1314	28	22362	12	34.724557
4	5	100724	972	18	19396	12	32.364317

	LONGITUDE	ADM_RATE_ALL	SIZE	...	MEDIAN_PLUS_LN_DEBT_GRADS	\
0	-86.568502	0.8986	4990	...	16106.0	
1	-86.799345	0.9211	13186	...	16954.0	
2	-86.174010	1.0000	351	...	NaN	
3	-86.640449	0.8087	7458	...	16550.0	
4	-86.295677	0.9774	3903	...	18952.0	

	Y2_LN_DEFAULT_RT	Y2_LN_DELINQ_RT	FED_UG_PAIDINFULL	\
0	0.172640	0.074776	0.007926	
1	0.060389	0.034483	0.059675	
2	0.157191	0.050167	NaN	
3	0.055035	0.031390	0.092947	
4	0.173804	0.061713	0.007872	

	UG_GRAD_2Y_FORBEAR_RT	LPSTAFFORD_CNT	NUM_STU_PLUS_LN_BAL	LPPPLUS_AMT	\
0	0.266409	31374.0	5201	113949554.0	
1	0.156121	56997.0	3727	87421879.0	
2	0.000000	4463.0	16	309618.0	
3	0.100349	19702.0	1397	29789762.0	
4	0.286976	34246.0	4602	95359062.0	

	FED_SCHOOL_CD	CLUSTERLABEL
0	001002	21
1	001052	5
2	016885	10
3	001055	5
4	001005	21

[5 rows x 50 columns]

[18]: *# Function returns cluister label given a school id*

```
def get_label(unit_id):

    rslt_df = df[df['UNITID'] == int(unit_id)]
    cluster_label = rslt_df.CLUSTERLABEL.iloc[0]

    return cluster_label

# Test function
# print('CLUSTER LABEL is: ', get_label(100706))
```

[19]: *# Function returns school name given a unique school identifier*

```
def get_school_name(unit_id):
```

```

rslt_df = df[df['UNITID'] == int(unit_id)]

school_name = rslt_df.SCHOOL_NAME.iloc[0]

return school_name

# Test function
# print('SCHOOL NAME is: ', get_school_name(100706))

```

```

[20]: # Given a unique school id, return a list of recommended schools in the same_
      ↪ cluster
# Note: To find recommendations for Penn State, use UNITID = 214777
unit_id = input('What is the school id of a college which interests you? ')

# Get cluster label for selected school
cluster_label = get_label(int(unit_id))

# Return school_names for all schools in the cluster
rslt_df = df[df['CLUSTERLABEL'] == cluster_label]

# Print output of recommendations
pd.options.display.max_rows = 3000
print('Since you like', get_school_name(unit_id), ', then you might also be_
      ↪ interested in these other schools.')
rslt_df.SCHOOL_NAME

```

What is the school id of a college which interests you? 214777  
 Since you like Pennsylvania State University-Main Campus , then you might also  
 be interested in these other schools.

```

[20]: 1308    Pennsylvania State University-Penn State Erie-...
      1309    Pennsylvania State University-Penn State New K...
      1310    Pennsylvania State University-Penn State Shenango
      1311    Pennsylvania State University-Penn State Wilke...
      1312    Pennsylvania State University-Penn State Scranton
      1313    Pennsylvania State University-Penn State Lehigh...
      1314    Pennsylvania State University-Penn State Altoona
      1315    Pennsylvania State University-Penn State Beaver
      1316    Pennsylvania State University-Penn State Berks
      1317    Pennsylvania State University-Penn State Harri...
      1318    Pennsylvania State University-Penn State Brand...
      1319    Pennsylvania State University-Penn State Fayette...
      1320    Pennsylvania State University-Penn State Hazleton
      1321    Pennsylvania State University-Main Campus
      1322    Pennsylvania State University-Penn State Great...

```

```
1323    Pennsylvania State University-Penn State Mont ...
1324    Pennsylvania State University-Penn State Abington
1325    Pennsylvania State University-Penn State Schuy...
1326        Pennsylvania State University-Penn State York
1770                Arizona State University-West
1779                Arizona State University-Polytechnic
1856                Arizona State University-Downtown Phoenix
1944                Pennsylvania State University-World Campus
1984                Arizona State University-Skysong
Name: SCHOOL_NAME, dtype: object
```

```
[ ]:
```